

# Concurrency (I)

Lecture 12 (18<sup>th</sup> May 2021)

-- *Liang: Web Chapter 32* --

# What is concurrency?



- Performing operations concurrently (in parallel)
  - We can walk, talk, breathe, see, hear, smell... all at the same time;
  - Computers can do this as well: download a file, print a file, receive email, run the clock, more or less in parallel....
- How are these tasks typically accomplished?

# When to apply concurrency?

Three reasons for applying concurrency:

1. Responsiveness
2. Efficiency
3. (Simulation of) naturally concurrent processes

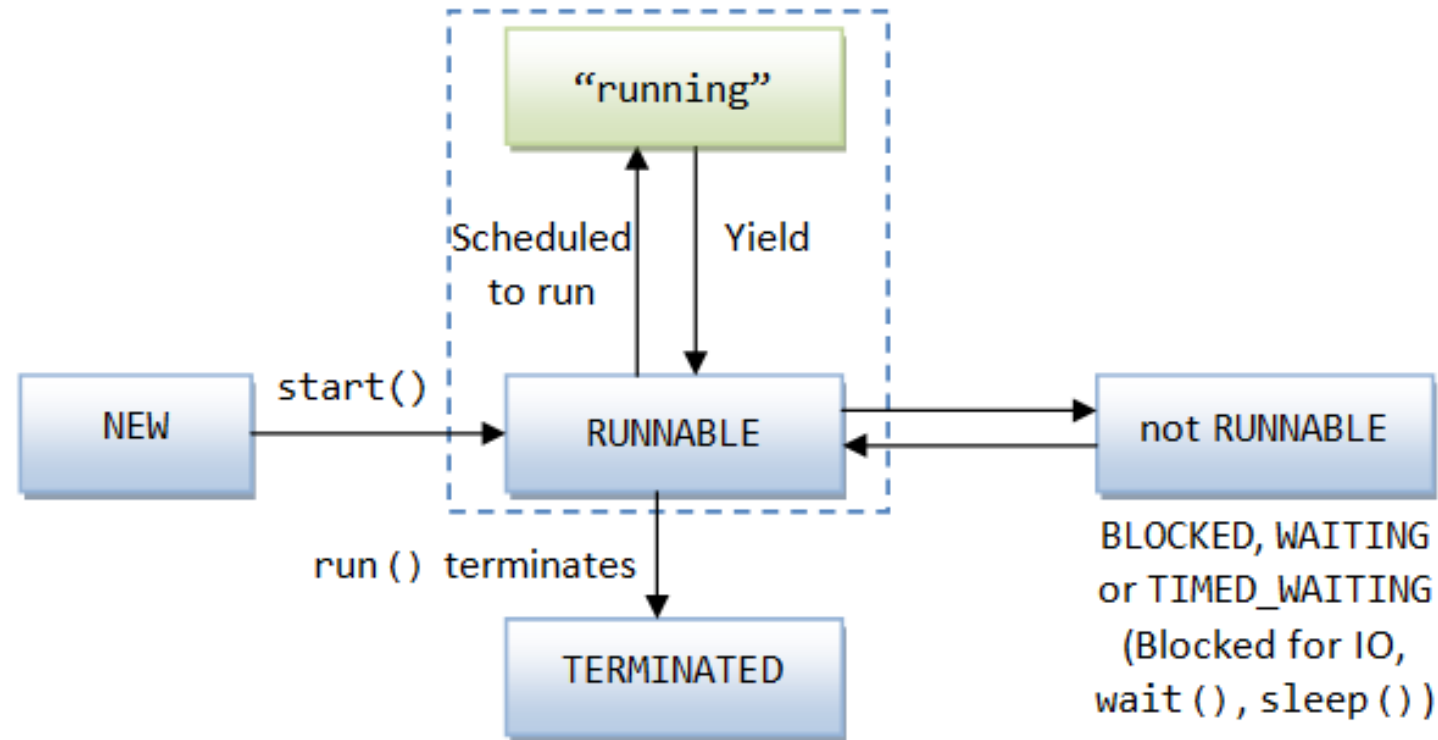
Furthermore:

When the order of actions is not relevant to the correctness of the program, concurrency enables to avoid specifying this order, resulting in simpler program code.

# About threads

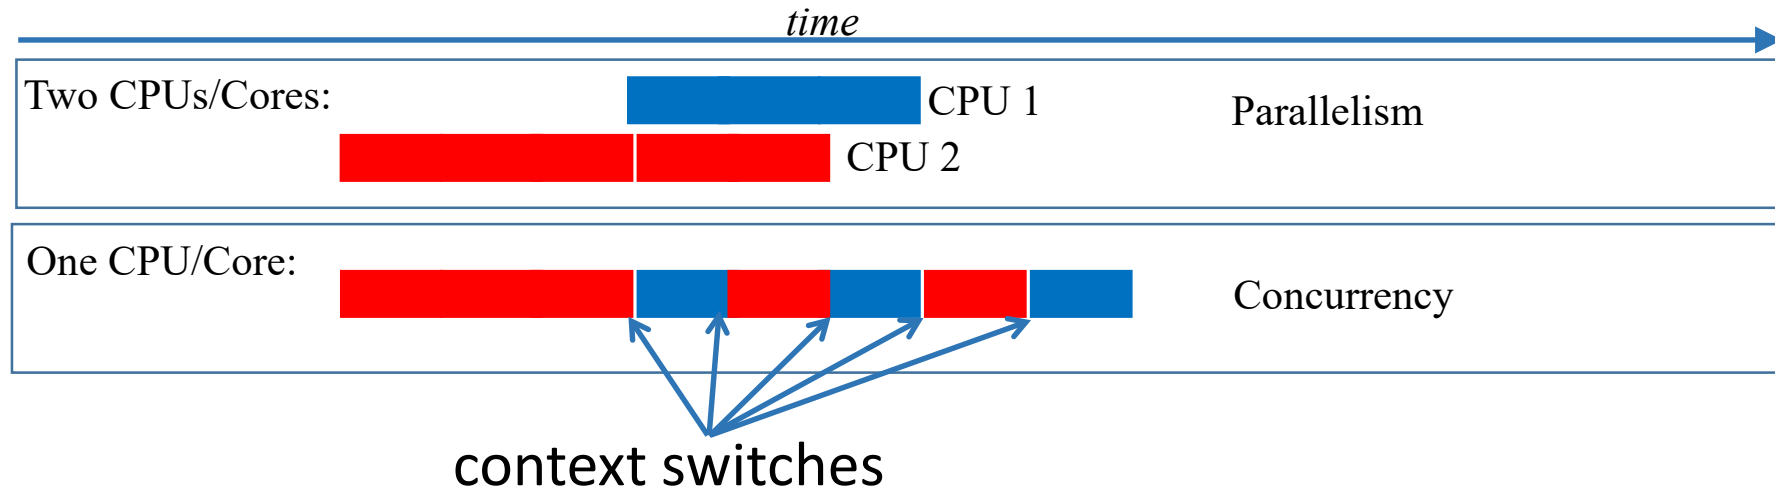
- The program *flow of control* is the order in which the computer executes the statements in a program.
- A **thread** corresponds to a *single sequential flow of control* (within a concurrent program)
- A thread executes independently of others, while at the same time sharing underlying system resources and objects constructed within the same program
- Each thread has its own call stack: threads do *not* share local variables and parameters
- Every program consists of at least one thread.

# States and state transitions of threads



- The operating system manages the threads. A so-called *scheduler* controls the transitions. In Java this is done by the JVM (Java Virtual Machine)

# Two threads



## Terminology:

- **Time-slicing (aka interleaving)** : Threads *alternatingly* get a small amount of execution time on the CPU
- **Scheduler**: decides which thread, when, and for how long will get the CPU
- **Context switch**: When the scheduler gives another thread the CPU the state of the current thread is stored and the state of the other one is loaded. This is a relatively costly operation.

# Nondeterminism

- Concurrent programs are nondeterministic
  - *Different executions with the same input can give different results*
  - *E.g., which thread terminates first is often unpredictable*
- Concurrent programs are complex
  - Several threads are executed simultaneously
  - In each execution, threads can get different amounts of CPU-time. Unexpected values of variables and of results can occur (and should be avoided)
    - Context switches can occur halfway between variable updates
  - Statements can be executed in ‘surprising’ order (different than in the sequential case, this now may become visible)
- We come back to this later (Problems with threads)

*“... Threads discard the most essential and appealing properties of sequential computation: comprehensibility, predictability, and determinism. ...”*

*(Edward.A. Lee, The problems with Threads, 2006)*

# Java: Objects and Threads

- OOP lecture 1: What are objects?
- Answer: Building blocks of software systems
  - *a program is a collection of interacting objects*
- Objects have **capabilities** that allow them to perform specific actions
- Thread **capability**
  - a thread is able to **perform a task**.
- A **Task** is an object that provides a **run** method.
  - If a thread is asked to perform a task it will execute the task's run method
- In Java, programs can have multiple threads



# Creating a new thread(I): interface Runnable

1. Introduce a class implementing the interface **Runnable**. This class is said to be **active**
2. Implement the method **run** with the code for the concurrent task:

```
public void run() { // code of the task goes here }
```

Note: method **run**


- has no parameters
- no return value,
- cannot throw checked exceptions

```
interface Runnable {  
    void run();  
}
```

3. Create an instance of class **Thread** with an active class object as argument
4. Call on this instance of Thread the method **start**; method **run** will then be executed

```
public Thread( Runnable task )
```

# Creating a new thread(II): sequential example



```
public class Main {  
    public static void main(String[] args) {  
        Simulation sim = new Simulation();  
        sim.simulate();  
    }  
}
```

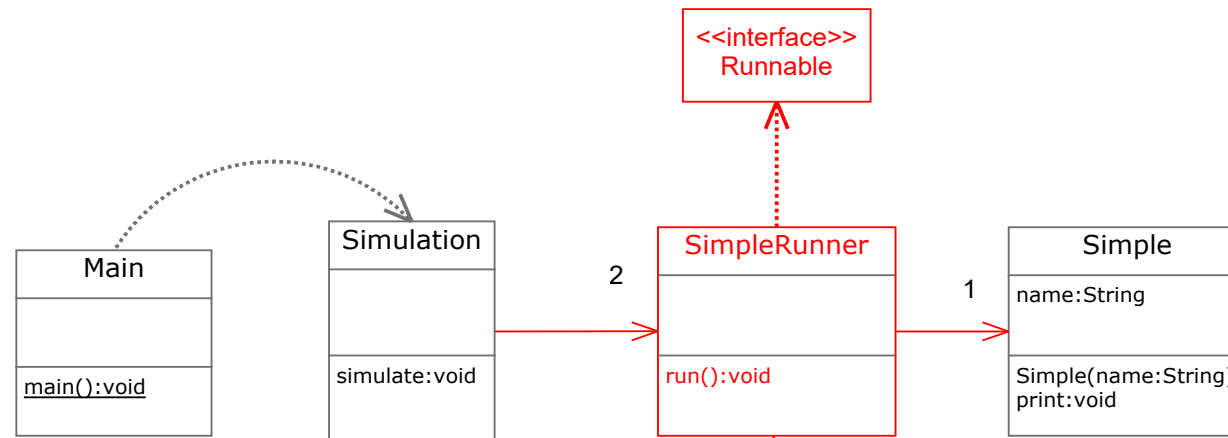
```
public class Simple {  
    private String name = null;  
  
    public Simple(String name) {  
        this.name = name;  
    }  
  
    public void print() {  
        System.out.println("My name =" + name);  
    }  
}
```

```
public class Simulation {  
    private static final int N = 3;  
  
    public void simulate() {  
        Simple s1 = new Simple("One");  
        Simple s2 = new Simple("Two");  
        for (int i = 0; i < N; i++) {  
            s1.print();  
        }  
        for (int i = 0; i < N; i++) {  
            s2.print();  
        }  
    }  
}
```

run:

```
My name =One  
My name =One  
My name =One  
My name =Two  
My name =Two  
My name =Two
```

# Creating a new thread (III): changed to concurrent example



```
run:
My name =Two
My name =One
My name =Two
My name =Two
My name =One
My name =One
```

```
public class SimpleRunner implements Runnable{
    private Simple simple = null;

    public SimpleRunner(Simple simple){
        this.simple = simple;
    }

    @Override
    public void run() {
        for (int i = 0; i < Simulation.N; i++) {
            simple.print();
        }
    }
}
```

```
public class Simulation {

    public static final int N = 3;

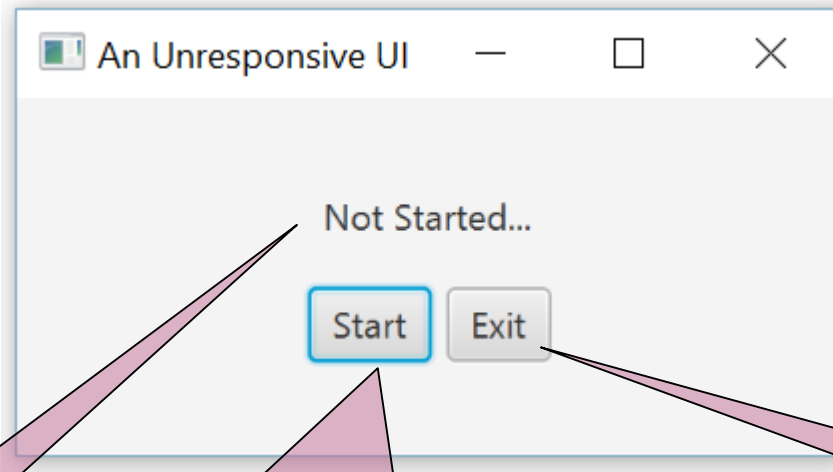
    public void simulate() {
        Simple s1 = new Simple("One");
        Simple s2 = new Simple("Two");
        SimpleRunner sr1 = new SimpleRunner(s1);
        SimpleRunner sr2 = new SimpleRunner(s2);
        Thread t1 = new Thread(sr1);
        Thread t2 = new Thread(sr2);
        t1.start();
        t2.start();
    }
}
```

```
public class Simple {
    private String name = null;

    public Simple(String name) {
        this.name = name;
    }

    public void print() {
        System.out.println("My name =" + name);
    }
}
```

# Nonresponsiveness (I)



Status label

Start button: starts a task for 10 seconds



Exit button: stops the program

```
Label statusLbl = new Label( "Not Started..." );  
Button startBtn = new Button( "Start" );  
Button exitBtn  = new Button( "Exit" );
```

```
public void start( Stage stage ) {  
    startBtn.setOnAction( e -> runTask() );  
    exitBtn.setOnAction( e -> stage.close() );  
}
```

# Nonresponsiveness (II)

```
startBtn.setOnAction( e -> runTask() );
```

```
public void runTask() {  
    for ( int i = 1; i <= 10; i++ ) {  
        String status = "Processing " + i + " of " + 10;  
        statusLbl.setText( status );  
        System.out.println( status );  
        takeABreak( 1000 );  
    }  
}
```

```
public void takeABreak( int millisec ) {  
    try {  
        Thread.sleep( millisec );  
    } catch ( InterruptedException e ) {  
        e.printStackTrace();  
    }  
}
```

# Thread.sleep



- `Thread.sleep`: static method in class `Thread`.
- A thread calling this method will be paused for the number of milliseconds given as an argument.
  - Note that it may be invoked in an ordinary (single-threaded) program to insert a pause in the single thread of that program
- It may throw a (checked) exception: `InterruptedException`
- Use:
  - To slow down programs, for, e.g., animations
  - Side effect: context switch

# TimeUnit: sleep

## Enum Constants

### Enum Constant and Description

DAYS

HOURS

MICROSECONDS

MILLISECONDS

MINUTES

NANOSECONDS

SECONDS

```
public void takeABreak( int millisec ) {  
    try {  
        TimeUnit.MILLISECONDS.sleep(millisec);  
    } catch ( InterruptedException e ) {  
        e.printStackTrace();  
    }  
}
```

# Analyzing the nonresponsive Program

- Two issues:
  1. Clicking Exit button has no effect until the task finishes
  2. Label is not updated
- Reason:
  - all UI event handlers in JavaFX run on a single thread:  
the **JavaFX Application Thread**
  - as long as one handler is executed, other handlers have to wait. In this case:
    - the handler for the exit button
    - the handler for updating the (view of the) label.



# Fixing a Nonresponsive Program Using Threads

- This is how to fix the problem:



- Have the button event handler method create a new thread to execute runTask
- Once created, the new thread will be an **independent process** that proceeds on its own
- Now, the work of the event handler is ended, and the main thread (that executed the handler) is ready to respond to something else
- If the “Exit” button is clicked while the new thread executes the task, then the program will end

## Multithreaded Program that tries to fix the Nonresponsive GUI

1. Introduce an (active) class **Task** implementing `run` which executes the `runTask` method.
2. Let the handler call a new method `startTask` which creates an instance of `Task`, and an instance of `Thread` with the former instance as argument.
3. Start the new thread, which now can run concurrently with the JavaFX Application Thread.

## Step 1: active class

```
public class Task implements Runnable {  
    ResponsiveGUI gui;  
  
    public Task ( ResponsiveGUI gui ) {  
        this.gui = gui;  
    }  
  
    public void run(){  
        gui.runTask();  
    }  
}
```

## Step 2 and 3: Create and start a thread

```
startBtn.setOnAction( e -> startTask() );
```

```
public void startTask() {  
    Task task = new Task( this );  
    Thread taskThread = new Thread( task );  
    taskThread.start();  
}
```

## Step 4: Run your program

After pressing the Start button you get:

```
Exception in thread "Thread-4" java.lang.IllegalStateException: Not on FX  
application thread; currentThread = Thread-4
```

In principle, the solution is ok. However, when running JavaFX (or Swing) applications you have to be more careful.

We will analyse the problem, and fix the error later.  
Let's first have a look at some other issues.

# Executing runnables (I)

```
public class PrintTask implements Runnable
{
    private final int sleepTime;
    private final String taskName;
    private final static Random generator = new Random();

    public PrintTask ( String name ) {
        taskName = name;
        sleepTime = generator.nextInt( 5000 );
    }

    public void run () {
        ... // next slide
    }
}
```

Implement **Runnable** to  
define a task that can  
execute concurrently

# Executing runnables (II)

Define task in method **run**

```
public void run() {  
    System.out.printf ( "%s going to sleep for %d ms.\n",  
                        taskName, sleepTime );  
    takeABreak( sleepTime );  
    System.out.printf ( "%s done sleeping\n", taskName );  
}
```

# Executing runnables (III)

```
public class TaskExecutor {  
    private final List<Runnable> tasks;  
  
    public TaskExecutor(List<Runnable> runnables) {  
        tasks = runnables;  
    }  
  
    public void executeConcurrently( ) {  
        System.out.println("Starting Tasks");  
        tasks.forEach(task -> new Thread(task).start());  
        System.out.println("Tasks started");  
    }  
}
```



# Executing runnables (IV)

```
public class Main {  
    private static final int NR_PRINT_TASKS = 3;  
  
    public static void main(String[] args) {  
        List<Runnable> tasks = IntStream.rangeClosed(1, NR_PRINT_TASKS)  
            .mapToObj(taskNr -> new PrintTask("Task " + taskNr))  
            .collect(Collectors.toList());  
        TaskExecutor te = new TaskExecutor(tasks);  
        te.executeConcurrently();  
        System.out.println("main ends");  
    }  
}
```

Starting Tasks

Tasks started

Task 1 going to sleep for 100 ms.

Task 2 going to sleep for 2462 ms.

Task 3 going to sleep for 1311 ms.

main ends

Task 1 done sleeping

Task 3 done sleeping

Task 2 done sleeping

# Executing runnables (V)

- Sometimes you want to test your program in a more predictable way: Do It Yourself

```
public class TaskExecutor {  
    List<Runnable> tasks;  
  
    public void executeSequentially( ) {  
        tasks.forEach(Runnable::run);  
    }  
}
```

# Executing runnables (VI)

```
public class Main {  
    private static final int NR_PRINT_TASKS = 3;  
  
    public static void main(String[] args) {  
        List<Runnable> tasks = IntStream.rangeClosed(1, NR_PRINT_TASKS)  
            .mapToObj(taskNr -> new PrintTask("Task " + taskNr))  
            .collect(Collectors.toList());  
        TaskExecutor te = new TaskExecutor(tasks);  
        te.executeSequentially();  
        System.out.println("main ends");  
    }  
}
```

```
Task 1 going to sleep for 2727 ms.  
Task 1 done sleeping  
Task 2 going to sleep for 1890 ms.  
Task 2 done sleeping  
Task 3 going to sleep for 1601 ms.  
Task 3 done sleeping  
main ends
```

# Sharing resources



# Atomicity and scheduling

A (Java) statement such as `i++;` is not executed as one (indivisible = atomic) step, but consists of a sequence of smaller (assembly) steps.

Example for `i++;` (depends on processor type):

```
LOAD @i, r0;    load the value of 'i' into a register from memory
ADD r0, 1;      increment the value in the register
STORE r0, @i;   write the updated value back to memory
```

**Scheduling is done on the level of assembly code** (or even lower) and hence during the execution of `i++;` a context switch can occur



# Two possible scenarios

Threads T1 and T2 update a common variable *i*

**Scenario as you might expect**

T1	T2	R/W	i
			0
Read		←	0
+ 1			0
Write		→	1
	Read	←	1
	+ 1		1
	Write	→	2

**Scenario with unexpected result**

T1	T2	R/W	i
			0
Read		←	0
	Read	←	0
+ 1			0
	+ 1		0
Write		→	1
	Write	→	1

- When *i++* is executed by two different threads, the end result is that *i* can be both 2 and 1 larger than at the start
- NB: this is irrespective of whether the threads are executed on one or two processors

# Unsynchronized Data Sharing Demo (I)

```
public class Buffer // CAUTION: NOT THREAD SAFE!
{
    private final int[] bufferArray;
    private int writeIndex = 0;

    public static final int PAUSE = 50;

    public Buffer( int size ) {
        bufferArray = new int[ size ];
    }

    public void add( int value ) {
        ... // next slide
    }
}
```

## Unsynchronized Data Sharing Demo (II)

```
public void add( int value ) {  
    int position = writeIndex;  
  
    takeABreak();  
  
    bufferArray[ position ] = value;  
  
    takeABreak();  
  
    writeIndex = position + 1;  
}
```





# Unsynchronized Data Sharing Demo (III)

```
public class BufferWriterTask implements Runnable
{
    private final Buffer myBuffer;
    private final int startValue, nrOfItems;

    public BufferWriterTask( int start, int nr_of_items, Buffer array ) {
        startValue = start;
        nrOfItems   = nr_of_items;
        myBuffer     = array;
    }

    public void run() {
        for ( int i = 0; i < nrOfItems; i++ ) {
            myBuffer.add( startValue + i );
        }
    }
}
```

# Unsynchronized Data Sharing Demo (IV)

```
private static final int BSIZE = 10;

public static void main( String[] arg ) {
    Buffer buffer = new Buffer( BSIZE );

    BufferWriterTask writer1 = new BufferWriterTask( 1, BSIZE/2 , buffer );
    BufferWriterTask writer2 = new BufferWriterTask( 14, BSIZE/2 , buffer );

    Thread writer_thread1 = new Thread ( writer1 );
    Thread writer_thread2 = new Thread ( writer2 );

    writer_thread1.start();
    writer_thread2.start();

    System.out.println( buffer );
}
```



# Unsynchronized Data Sharing Demo (V)

You can use the **join** method to force one thread to wait for another thread to finish.

```
public static void main( String[] arg ) {  
  
    <....>  
    // wait for the completion of the other threads  
    try {  
        writer_thread1.join();  
        writer_thread2.join();  
    } catch ( InterruptedException ex ) {  
        ex.printStackTrace();  
    }  
  
    // display the contents of the shared array  
    System.out.println( buffer );  
}
```



# Race conditions

- A **race condition**: possibly incorrect result caused by the update of a common variable/object by several threads.

```
public void add( int value ) {  
    int position = writeIndex;  
  
    takeABreak();  
  
    bufferArray[ position ] = value;  
  
    takeABreak();  
  
    writeIndex = position + 1;  
}
```

# Non-atomic Operations

- Writing

```
public void add( int value ) {  
    buffer[ writeIndex++ ] = value;  
}
```

won't work: `array[ writeIndex++ ] = value;` is not atomic (single unit of work, which cannot be interrupted in the middle).

- Without special measures, atomicity in Java is hard to obtain: almost any Java statement consists of several non-atomic operations.

# Thread Synchronization (I)

- Coordinates access to shared data by multiple concurrent threads
  - Indeterminate results may occur unless access to a shared object is managed properly
  - *Code* (typically the body of a method) that manipulates shared data is called a **critical section**.
- Basic idea: Give only one thread at a time exclusive access to a critical section.
- Other threads have to wait
- When the thread with exclusive access leaves the critical section (i.e. finishes manipulating the shared object), one of the threads that was waiting is allowed to proceed



# Thread Synchronization (II)

- Exclusive access to such a critical section (also called **mutual exclusion**) is obtained through **locking**.
- The basic form of locking in Java is by means of so-called
  1. *synchronized methods*
  2. *synchronized blocks*
- Any Java object can act as a lock.
  - when calling a synchronized method, the current object (**this**) is used for locking.
  - in a synchronized block the lock object is specified explicitly.
- When a thread calls a synchronized method (or enters a synchronized block) the object's lock, when it is free, will become occupied (by the executing thread). When the lock is not free, the calling/entering thread has to wait.



# Thread Synchronization JAVA

- **synchronized** statement

- Enforces mutual exclusion on a block of code

```
synchronized ( object ) {  
    statements  
} // end synchronized statement
```

where *object* is the object whose lock will be acquired (normally **this**)

- A **synchronized** method is equivalent to a **synchronized** statement that encloses the entire body of a method, using the lock of **this**





# Synchronized Data Sharing—Making Operations Atomic (III)

```
public synchronized void add( int value ) {  
    int position = writeIndex;  
  
    bufferArray[ position ] = value;  
  
    writeIndex = position + 1;  
}
```


Using synchronized prevents more than one thread at a time from calling this method



# Visualization of race conditions (I)


- The cause of a race condition are *updates* (state changes) of an object that is *shared* between several threads
- Visualization: Use an **UML activity diagram** in which the shared object is explicitly indicated
- In case of accessing the shared object distinguish between read and write/update
- In the shared object *explicitly* indicate the *attribute* that is accessed

## Visualization of race conditions (II)



```
void inc() {  
    i++;  
}
```

Two threads use a shared object `x:X` with attribute `i` and try to update attribute `i`



**Check then act  
(or read-modify-write)**

Two threads use a shared object `x` with attribute `var`

# Avoiding race conditions

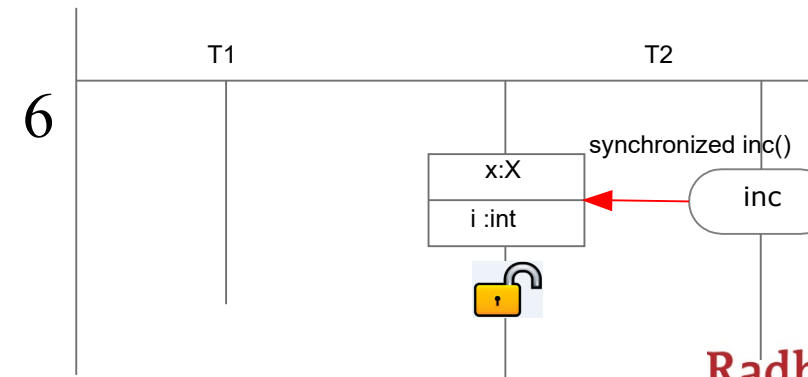


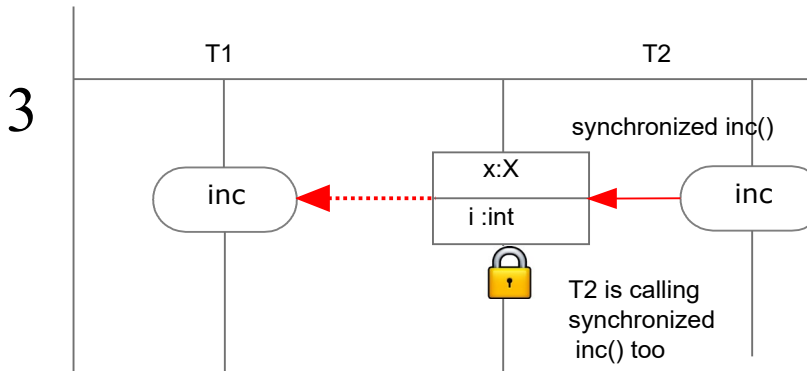
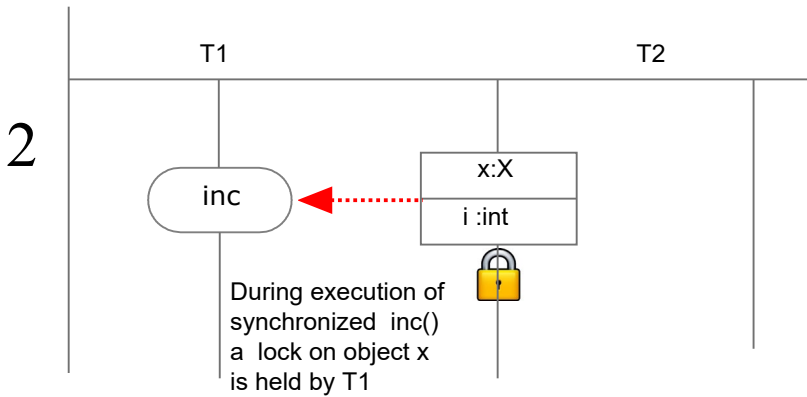
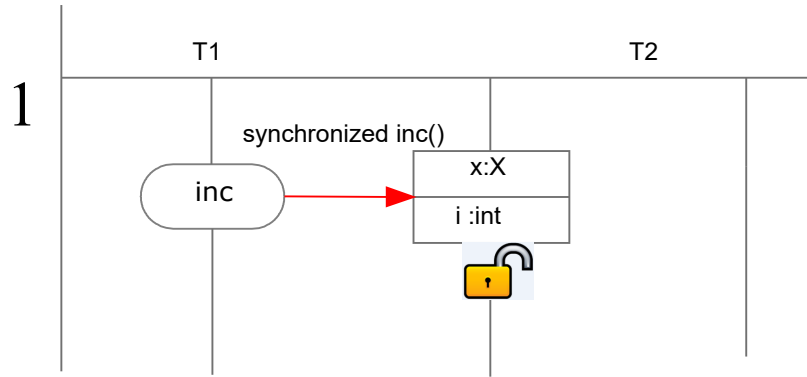
```
public synchronized void inc() {  
    i++;  
}
```

... or (equivalently) ...

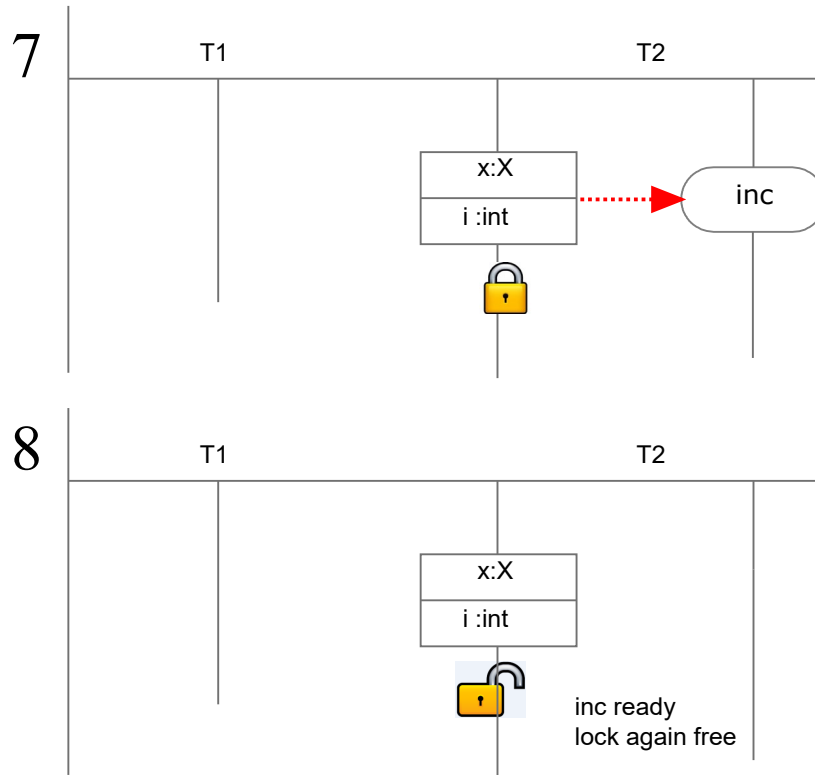
```
public void inc() {  
    synchronized (this) {  
        i++;  
    }  
}
```

- Method inc now can only be executed by one thread at the time.
- The use of synchronized will decrease performance, because the JVM has to do additional bookkeeping, so be frugal in the use of synchronized methods.

# Visualizing synchronized methods (I)



## Visualizing synchronized methods (II)



These figures are called a Lock Allocation Diagrams or LADs

# Check-Then-Act

- **Check-then-act**: if a context switch occurs *after* checking a condition but *before* the subsequent action is performed, the condition may have changed. This is dangerous if that action relies on the condition.


# CTE-example (I)

```
public class Counter {  
    private int counter;  
  
    public synchronized void incr() {  
        counter++;  
    }  
  
    public synchronized int getCounter() {  
        return counter;  
    }  
}
```



# CTE-example (II)

```
public class Incrementer implements Runnable {  
    private Counter counter;  
    private int limit;  
  
    public Incrementer( Counter counter, int limit ) {  
        this.counter = counter;  
        this.limit    = limit;  
    }  
  
    @Override  
    public void run() {  
        while ( counter.getValue() != limit ) {  
            counter.incr();  
        }  
    }  
}
```



# CTE-example (III)

```
public class CounterTest {
    public static void main( String[] args ) {
        Counter counter      = new Counter();
        Incrementer incr1    = new Incrementer( counter, 50 );
        Incrementer incr2    = new Incrementer( counter, 50 );

        Thread incr_thread1  = new Thread( incr1 );
        Thread incr_thread2  = new Thread( incr2 );

        incr_thread1.start();
        incr_thread2.start();

        try {
            incr_thread1.join();
            incr_thread2.join();
        } catch ( Exception e ) {
            e.printStackTrace();
        }
        System.out.println("Final counter: "+ counter.getValue());
    }
}
```

# CTE-example (IV)

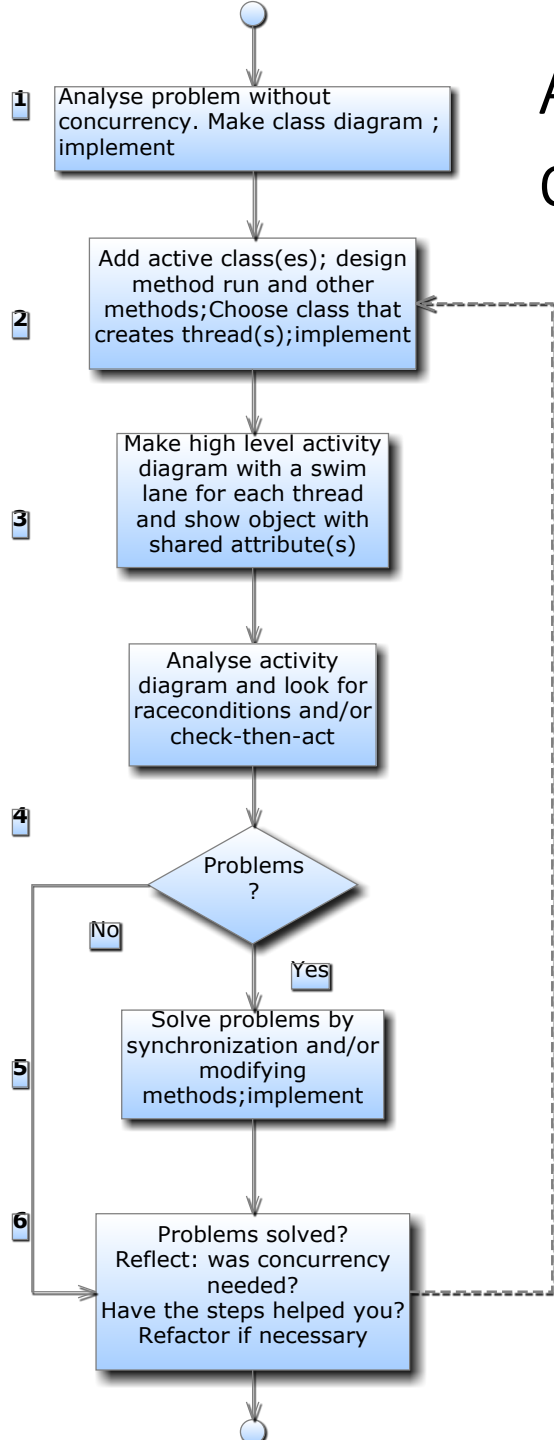
- Check-then-act problem.
- Synchronizing `getValue` and `incr` is insufficient
- Solution: put `getValue` and `incr` in a synchronized block

```
public class Incrementer implements Runnable {
    private Counter counter;
    private int limit;

    public Incrementer( Counter counter, int limit ) {
        this.counter = counter;
        this.limit    = limit;
    }

    @Override
    public void run() {
        boolean ready = false;
        while ( ! ready ) {
            synchronized ( counter ) {
                if ( counter.getValue() != limit ) {
                    counter.incr();
                } else {
                    ready = true;
                }
            }
        }
    }
}
```

# Approach for designing a concurrent program



- This approach consists of 6 steps and guides you to solve a problem by using concurrency, and to discover and handle possible issues caused by the use of threads.
- Step 1 and 2 involve creating a concurrent program;
- Step 3, 4 and 5 are necessary to solve the possible (synchronization) problems that are introduced by using threads
- Step 6 is a final check

