

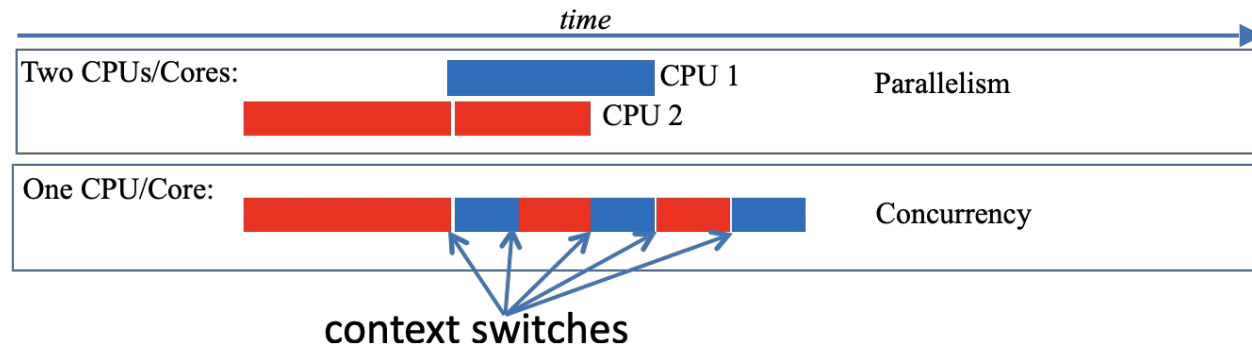
Concurrency: Threads

Tutorial 13 (19 May 2021)

Daniël Kuijsters


WE START AT 8:30

Parallelism Versus Concurrency

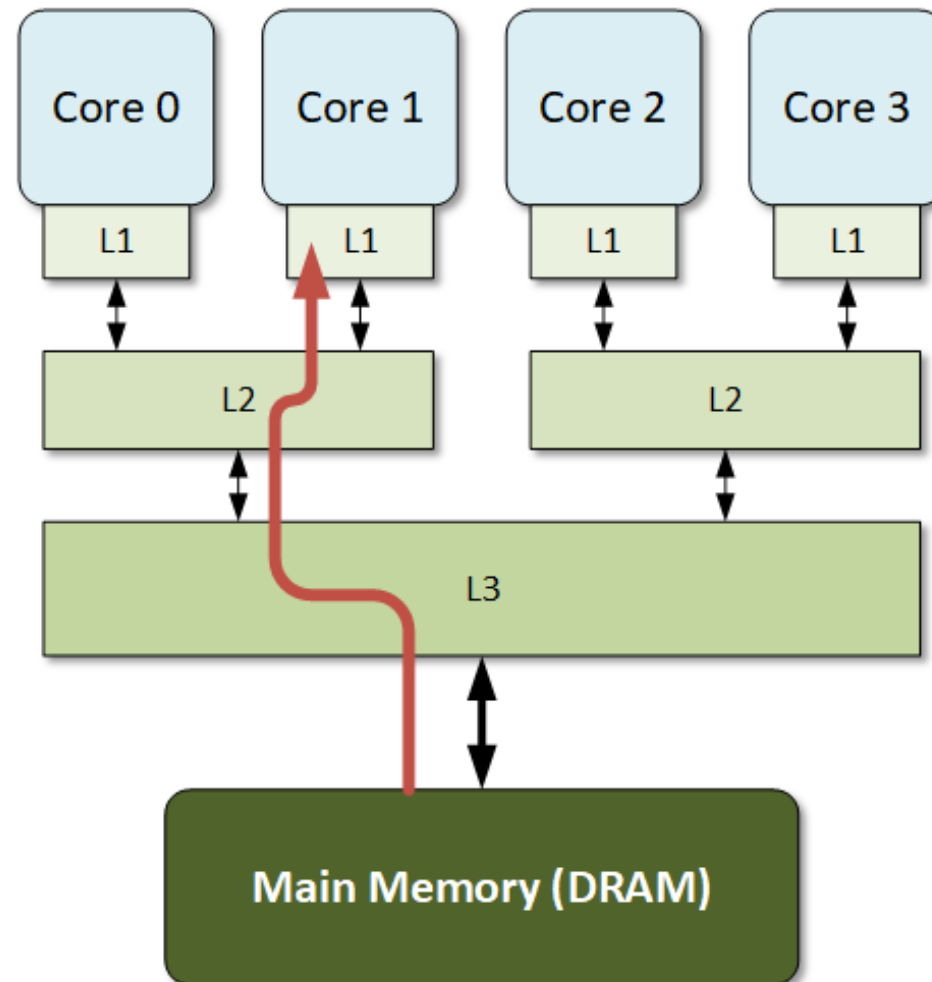


- Concurrency: At least two tasks are making progress at the same time. Used to hide latency.
- Parallelism: At least two tasks are being executed at the same time. Used to increase the speed or throughput.
- Parallelism implies concurrency.

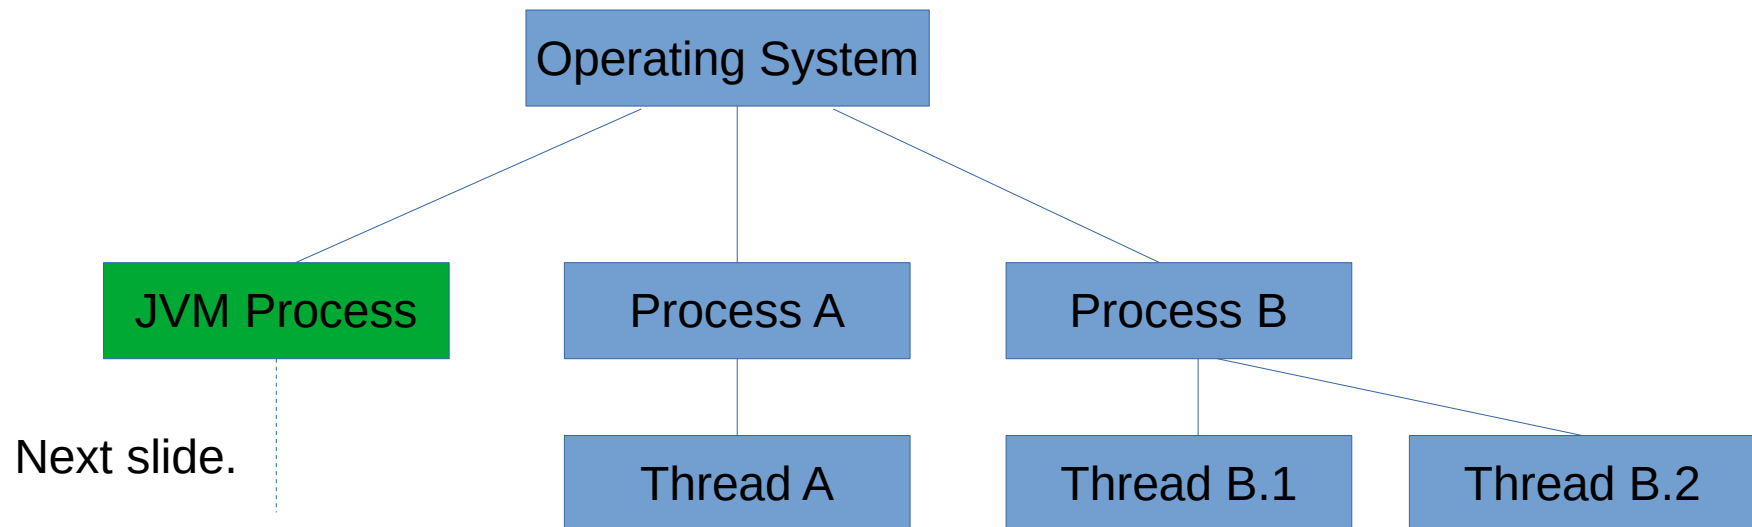
Moore's Law



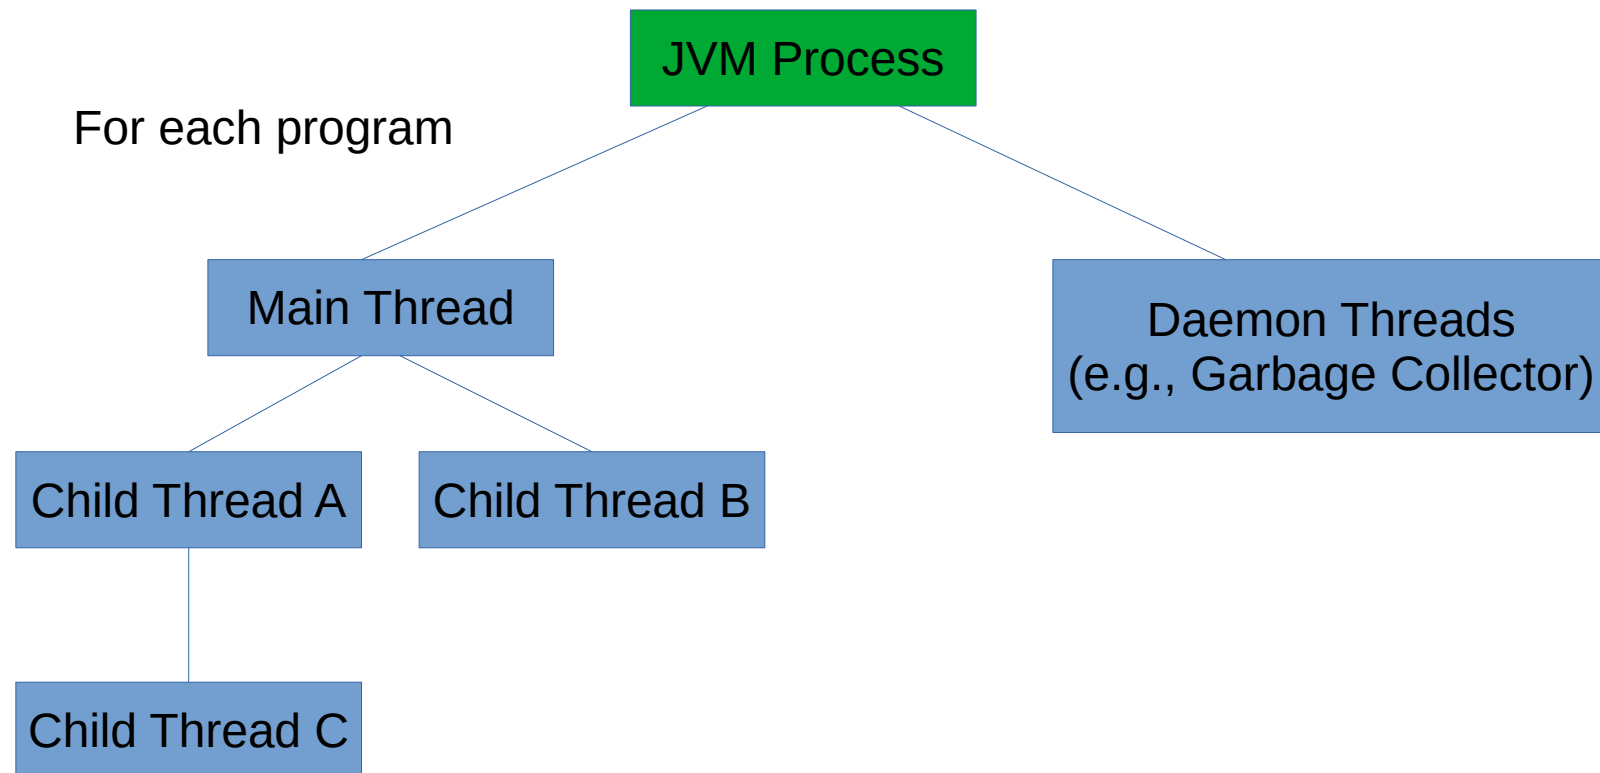
Possible Memory Architecture Multicore System



Java Virtual Machine (JVM) Process



JVM Threads



How Do You Create Threads?

- Implement the Runnable interface

```
@FunctionalInterface  
public interface Runnable {  
    public abstract void run();  
}
```

- Extend the Thread class

Extending Thread

```
public class Thread implements Runnable {  
    ...  
}  
  
public class Task extends Thread {  
    @Override  
    public void run() {  
        System.out.println("Run method executed by child thread.");  
    }  
  
    public static void main(String[] args) {  
        Task t = new Task();  
        t.start();  
        System.out.println("Main method executed by main thread.");  
    }  
}
```

DON'T DO THIS!

Implement Runnable: Using Classes

```
public class Task implements Runnable {  
    @Override  
    public void run() {  
        System.out.println("Run method executed by child thread.");  
    }  
  
    public static void main(String[] args) {  
        Task t = new Task();  
        Thread thread = new Thread(t);  
        thread.start();  
        System.out.println("Main method executed by main thread.");  
    }  
}
```

Implement Runnable: Using Classes

```
public class Task extends SomeSuperClass implements Runnable {  
    @Override  
    public void run() {  
        System.out.println("Run method executed by child thread.");  
    }  
  
    public static void main(String[] args) {  
        Task t = new Task();  
        Thread thread = new Thread(t);  
        thread.start();  
        System.out.println("Main method executed by main thread.");  
    }  
}
```

Replacing start() with run(): What Happens?

```
public class Task implements Runnable {  
    public Task() {  
        Thread t = new Thread(this);  
        t.run();  
    }  
  
    @Override  
    public void run() {  
        System.out.println("test");  
    }  
  
    public static void main(String[] args) {  
        new Task();  
    }  
}
```

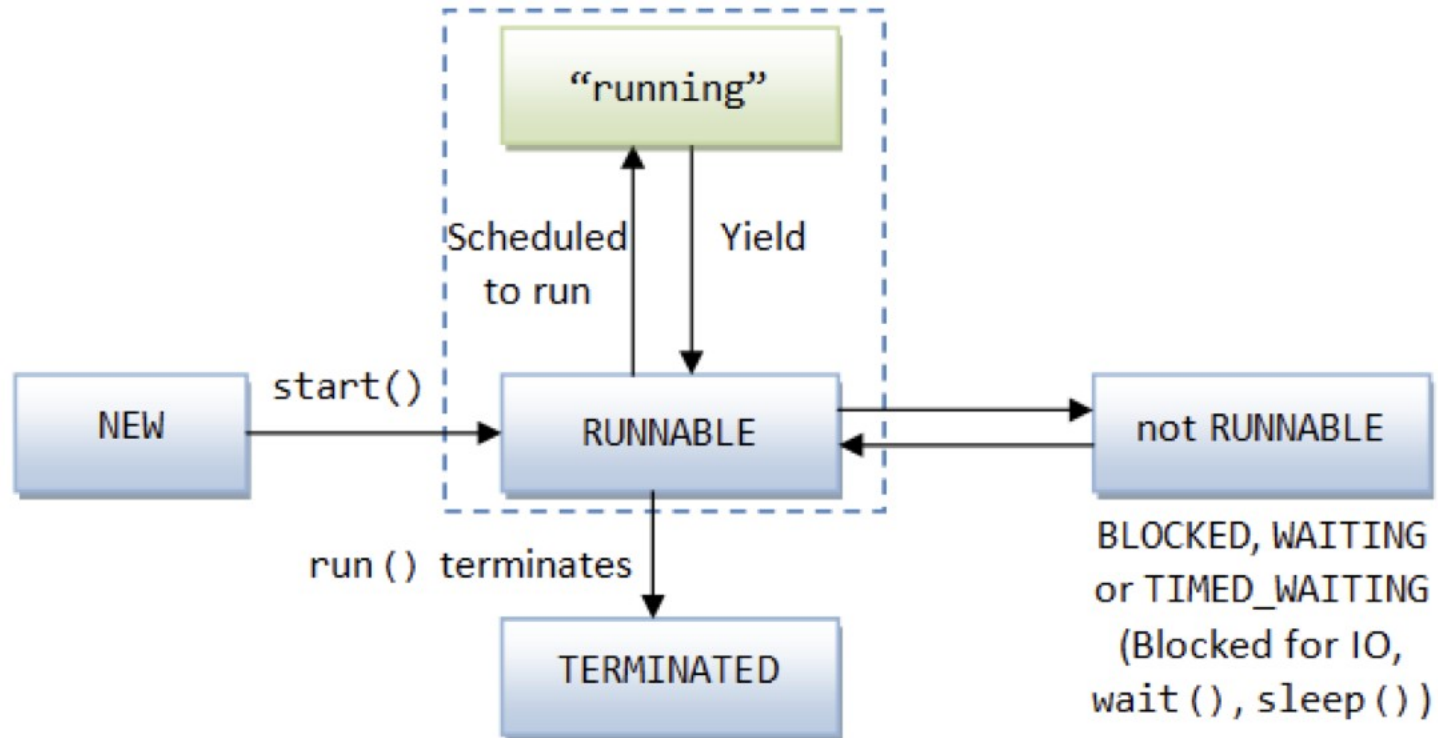
What Is Wrong In the Following Program?

```
public class Task implements Runnable {  
    public Task() {  
        Task t = new Task();  
        new Thread(t).start();  
    }  
  
    @Override  
    public void run() {  
        System.out.println("test");  
    }  
  
    public static void main(String[] args) {  
        new Task();  
    }  
}
```

What Is Wrong In the Following Program (2)?

```
public class Task implements Runnable {  
    public Task() {  
        Thread t = new Thread(this);  
        t.start();  
        t.start();  
    }  
  
    @Override  
    public void run() {  
        System.out.println("test");  
    }  
  
    public static void main(String[] args) {  
        new Task();  
    }  
}
```

Recall: Thread Finite State Machine



What Is Wrong In the Following Program (3)?

```
public class Main {  
    private class Task implements Runnable {  
        @Override  
        public void run() {  
            System.out.println("test");  
        }  
    }  
  
    public static void main(String[] args) {  
        Task task = new Task();  
        Thread thread = new Thread(task);  
        thread.start();  
    }  
}
```

Implement Runnable: Using Anonymous Inner Classes

```
public class Task {  
    public static void main(String[] args) {  
        Runnable r = new Runnable() {  
            @Override  
            public void run() {  
                System.out.println("Run method executed by child thread.");  
            }  
        };  
        Thread thread = new Thread(r);  
        thread.start();  
        System.out.println("Main method executed by main thread.");  
    }  
}
```



Implement Runnable: Using Lambda Expressions

```
public class Task {  
    public static void main(String[] args) {  
        Runnable r = () -> {  
            System.out.println("Run method executed by child thread.");  
        };  
        Thread thread = new Thread(r);  
        thread.start();  
        System.out.println("Main method executed by main thread.");  
    }  
}
```

Implement Runnable: Using Lambda Expressions

```
public class Task {  
    public static void main(String[] args) {  
        Thread thread = new Thread() -> {  
            System.out.println("Run method executed by child thread.");  
        };  
        thread.start();  
        System.out.println("Main method executed by main thread.");  
    }  
}
```

JVM Memory Model




Sharing Data Between Threads Is Difficult!

How to solve this data sharing issue?

- Do NOT share data at all → This week's first assignment.
- Wait for their execution to finish and combine data → This week's second assignment.
- Make use of synchronized statements and/or locks → The next weeks.

Independent Subproblems WITHOUT Combiner



Also called “embarrassingly parallel” problems.

How Many Threads Should I Create?

- Not an exact science; avoid “too little” and “too many”.
- For compute-intensive tasks: $\text{\#threads} = \text{\#cores} + 1$.
- For tasks that include I/O or other blocking operations: $\text{\#threads} = \text{\#cores} * (1 + \text{wait time} / \text{compute time})$.
- In other words: experiment!

```
int N_CPUS = Runtime.getRuntime().availableProcessors();
```

How Many Threads Should I Create (2)?

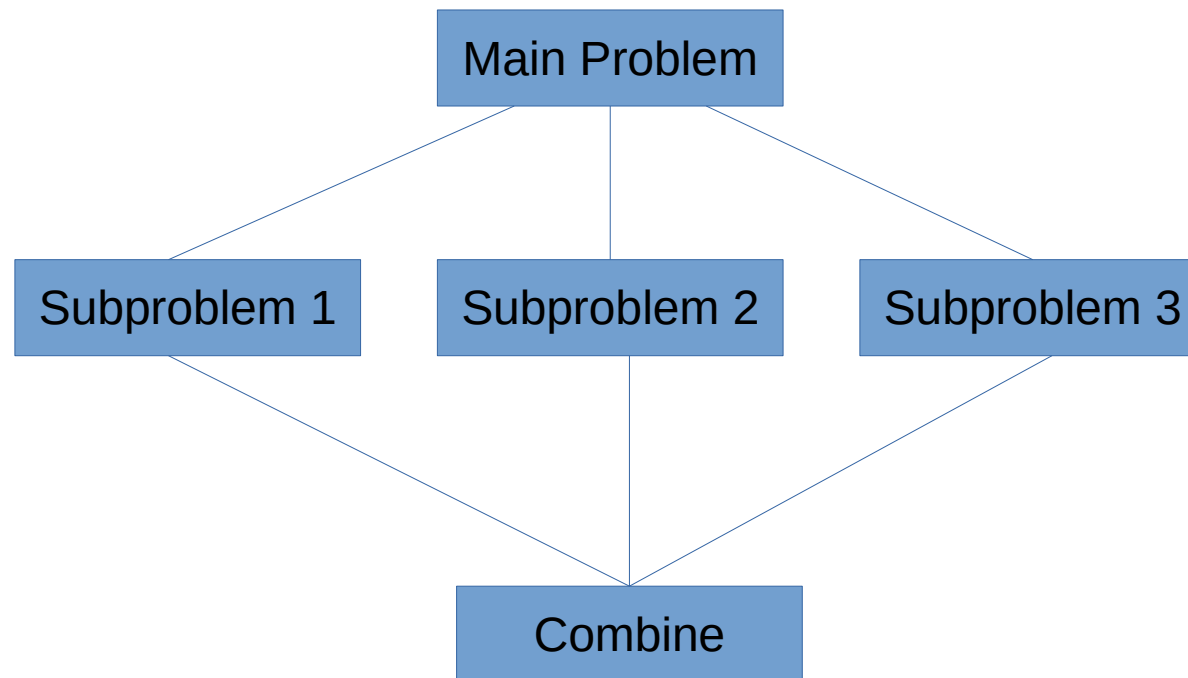
- Main thread is a thread itself, so let it do work too.
- Don't forget: creating a thread is NOT for free (For this reason, we will use thread pools starting next week).
- Typically we set a threshold below which we use a sequential solution.

Demo: Parallel Find Element


Assignment: FileFinder

- Sequential program that locates a given file in a directory.
- Since there can be many subdirectories and files, this procedure can be slow.
- Goal: speed this up using threads.
- Instantiation of “independent subproblems without combiner”.

Independent Subproblems WITH Combiner

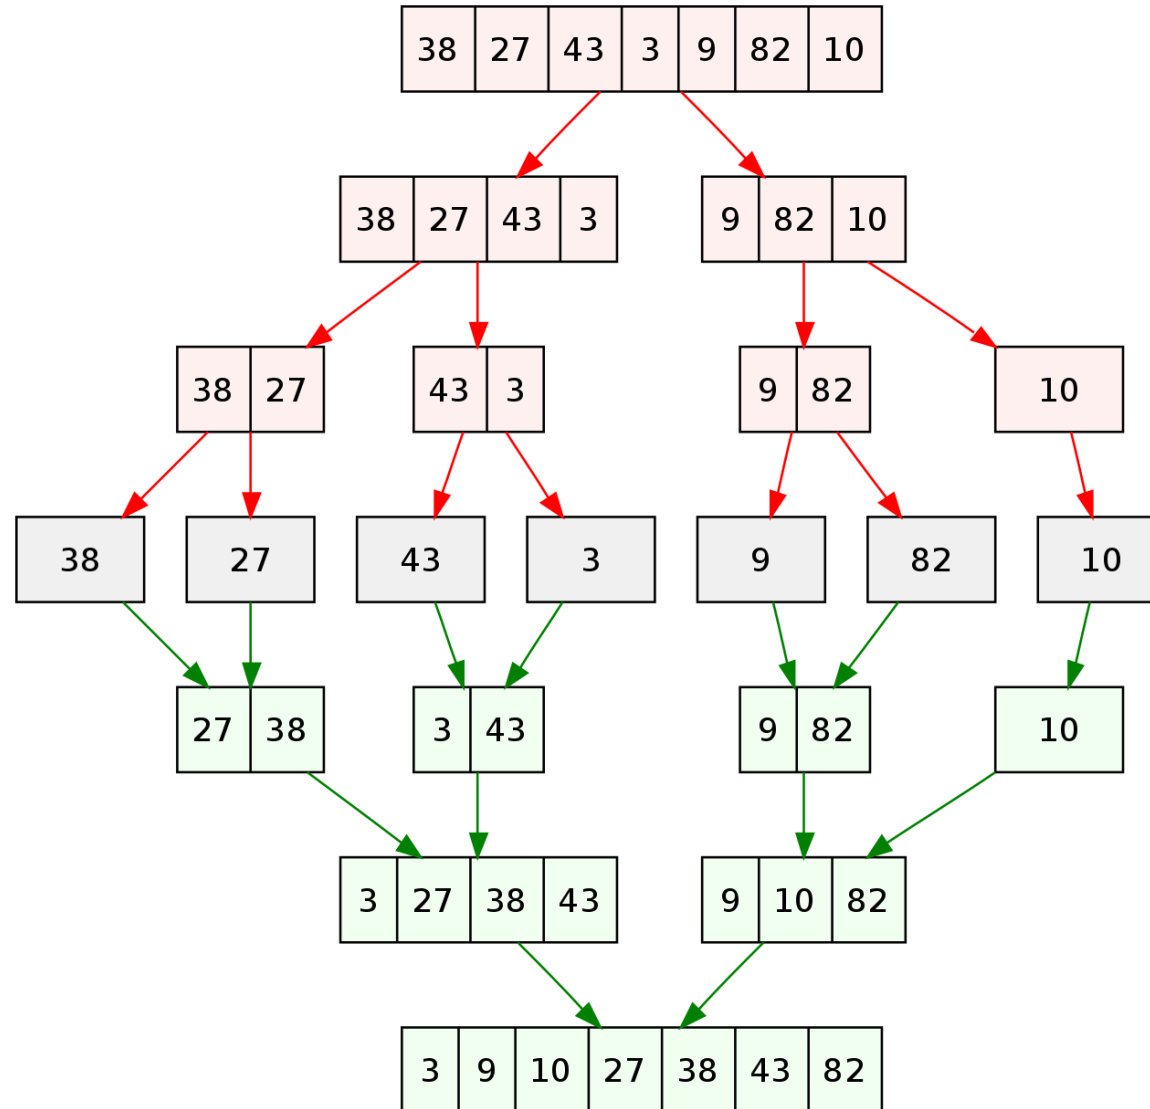


A.K.A. Fork-Join Model



Demo: Parallel Compute Max

Merge Sort



Assignment: Parallel MergeSort

- Merge sort is a sequential sorting algorithm.
- It is very suited for parallelization.
- Goal: speed this up using threads.
- Instantiation of “independent subproblems with combiner”.

